


Hidden Markov Models

Contents:

- Markov process, observable Markov models
- Hidden Markov models
- Problem 1: Scoring and evaluation
- Problem 2: Decoding
- Problem 3: Training

Markov model

- the current state of the system depends only on the previous state, not on the sequence before that
- the state of the system at time t is q_t
- the transition probability depends only on the previous state:



$$P(q_t = S_j | q_{t-1} = S_i, q_{t-2} = S_k \dots) = P(q_t = S_j | q_{t-1} = S_i)$$

$$a_{ij} = P(q_t = S_j | q_{t-1} = S_i)$$


Markov model

- stochastic model
- used to model a random system that changes state according to a transition rule that depends only on the current state

- Characterized by a set of N states

$$S = \{S_0, S_1, \dots, S_N\}$$

and transition probabilities from one state to another a_{ij}




Markov model properties

$$A = \begin{bmatrix} a_{00} & a_{01} \\ a_{10} & a_{11} \end{bmatrix}$$

$$a_{ij} \geq 0 \quad \forall i, j$$

$$\sum_{j=0}^N a_{ij} = 1 \quad \forall i$$



Stochastic matrix:


- each entry is non-negative
- rows sum up to 1

Example 1: Single fair coin observable process

- Observable: the output of the process is a set of states
- Outcomes:
 - Head – State 0
 - Tails – State 1

- Observed outcomes uniquely define a state sequence:
HHHTTTHHTT → 0001110011
- Transition probabilities:


$$A = \begin{bmatrix} 0.5 & 0.5 \\ 0.5 & 0.5 \end{bmatrix}$$



Example 2: Observable Markov model of weather

- Outcomes:
 - State 0 – Rainy
 - State 1 – Cloudy
 - State 2 – Sunny
- State transition probabilities:

$$A = \{a_{ij}\} = \begin{bmatrix} 0.4 & 0.3 & 0.3 \\ 0.2 & 0.6 & 0.2 \\ 0.1 & 0.1 & 0.8 \end{bmatrix}$$




Example 2: Observable Markov model of weather

- What is the probability that the weather for 8 consecutive days is sun, sun, sun, rain, rain, sun, cloudy, sun?
- Representing the information:
 - Observation sequence is:
 $O = \{\text{sun, sun, sun, rain, rain, sun, cloudy, sun}\}$
 - Corresponds to state sequence:
 $S = \{2, 2, 2, 0, 0, 2, 1, 2\}$
 - We need to calculate $P(O \mid \text{model})$
 $P(O \mid \text{model}) = P(S = \{2, 2, 2, 0, 0, 2, 1, 2\} \mid \text{model})$

Example 2: Observable Markov model of weather

$$A = \{a_{ij}\} = \begin{bmatrix} 0.4 & 0.3 & 0.3 \\ 0.2 & 0.6 & 0.2 \\ 0.1 & 0.1 & 0.8 \end{bmatrix}$$

π_i : initial state probability
 $\pi_i = P(q_1 = i)$




$$\begin{aligned} P(O \mid \text{model}) &= P(S = \{2, 2, 2, 0, 0, 2, 1, 2\} \mid \text{model}) \\ &= P(q_1 = 2)P(q_2 = 2 \mid q_1 = 2) \cdots P(q_8 = 2 \mid q_7 = 1) \\ &= \pi_2 \cdot a_{22} \cdot a_{22} \cdot a_{20} \cdot a_{00} \cdot a_{02} \cdot a_{21} \cdot a_{12} \\ &= \pi_2 \cdot (0.8)^2 (0.1) \cdot (0.4) \cdot (0.3) \cdot (0.1) \cdot (0.2) \end{aligned}$$

Hidden Markov models

- Observations are a probabilistic function of state
- The underlying sequence of states is not observable (it is hidden)
- Outputs are independent – observations are dependent only on the state that generated them, not on each other

Example: Two coins observable Markov process


- States: coins
- Observations:
 - Head
 - Tail
 - Each state can generate each observation with a certain probability
- The observed outcomes do not uniquely define state sequence
- Transition probabilities:



$$\begin{array}{ll} P(H) = P_1 & P(H) = P_2 \\ P(T) = 1 - P_1 & P(T) = 1 - P_2 \end{array}$$


$$A = \begin{bmatrix} a_{00} & 1 - a_{00} \\ 1 - a_{11} & a_{11} \end{bmatrix}$$

Example: urn and ball model




- 3 urns with 4 different color balls
- We do not see the urns, someone is extracting balls from them and tells us the color of each
- Steps:
 1. Select one urn at random
 2. Pick a ball from the urn, tell what color it is
 3. Put ball back to the urn
 4. Select new urn based on a random selection procedure from current urn
 5. Repeat steps 2-4

Example: urn and ball model



- Observations: the colors of the balls
- States: the identity of the urn
- State transitions: the selection process for next urn given current one

Example: urn and ball model




$$P(R)=0.20$$

$$P(G)=0.30$$

$$P(B)=0.10$$

$$P(Y)=0.40$$




$$P(R)=0.45$$

$$P(G)=0.15$$

$$P(B)=0.20$$

$$P(Y)=0.20$$



$$P(R)=0.15$$

$$P(G)=0.70$$

$$P(B)=0.10$$

$$P(Y)=0.05$$

- Urns contain different ratio of colours
- Observation sequence: R B Y Y G B Y G R ..
- The observation sequence (individual colors) do not reveal the state (which urn it comes from)

Discrete symbol observation HMM

- A set of N states

$$S = \{S_0, S_1, \dots, S_N\}$$

- Transition probabilities

$$a_{ij} = P(q_t = S_j | q_{t-1} = S_i)$$

- A set of M observation symbols


$$V = \{v_1, v_2, \dots, v_M\}$$

- Probability distribution (state j symbol k)

$$b_j(k) = P(o_t = v_k | q_t = j)$$

- Initial state distribution

$$\pi = \{\pi_i\} = P(q_1 = i)$$




Discrete symbol observation HMM


Specification of an HMM:

- Two model parameters, N and M
 - Number of states N
 - Number of symbols M
- Three probability measures A, B, π
 - Transition probability matrix A
 - State probability distribution B
 - Initial state distribution π

$$\lambda = (A, B, \pi)$$



Left-to-right HMM




$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

$$a_{ij} = 0 \quad j < i$$

$$\pi_i = \begin{cases} 0, & i \neq 1 \\ 1, & i = 1 \end{cases}$$

Ergodic HMM



$$a_{ij} > 0 \quad \forall i, \forall j$$

HMM as a symbol generator

An HMM with parameters N , M , A , B , and π can generate an observation sequence:

$$O = \{ o_1, o_2, o_3, \dots, o_T \}$$

1. Choose initial state $q_1 = i$ from the initial state distribution π
2. Set $t=1$
3. Choose $o_t = v_k$ according to distribution $b_j(k)$
4. Transition to state $q_{t+1} = j$ according to state transition probability a_{ij}
5. Set $t=t+1$
6. Repeat from step 3

HMM as a symbol generator

Time t	1	2	3	4	5	6	7	8	...	T
State	q_1	q_2	q_3	q_4	q_5	q_6	q_7	q_8	...	q_T
Observation	o_1	o_2	o_3	o_4	o_5	o_6	o_7	o_8	...	o_T

Think of the HMM as generating the observation sequence as it transitions from state to state

HMM problems

- Problem 1: Scoring and evaluation
 - How to compute efficiently the probability of an observation sequence O given the model λ ? (How to calculate $P(O|\lambda)$)
- Problem 2: Decoding
 - Given an observation sequence O and a model λ , how do we determine the corresponding state sequence q that best explains how the observations were generated?
- Problem 3: Training
 - How to adjust the parameters $\lambda=\{A,B, \pi\}$ to maximize the probability of generating a given observation sequence? (How to maximize $P(O|\lambda)$)

Problem 1: Scoring and evaluation

- Given an observation sequence $O = \{ o_1, o_2, o_3, \dots o_T \}$ we want to compute the probability of generating it $P(O|\lambda)$
- We assume a sequence of states $q = \{ q_1, q_2, q_3, \dots q_T \}$
- Decompose the problem by summing over all possible state sequences:

$$P(O|\lambda) = \sum_{\text{all } q} P(O|q, \lambda) P(q|\lambda)$$

Problem 1: Scoring and evaluation

- Given an observation sequence $O = \{ o_1, o_2, o_3, \dots o_T \}$ we want to compute the probability of generating it $P(O|\lambda)$
- We assume a sequence of states $q = \{ q_1, q_2, q_3, \dots q_T \}$
- Decompose the problem by summing over all possible state sequences:

$$P(O|\lambda) = \sum_{\text{all } q} P(O|q, \lambda) P(q|\lambda)$$

Likelihood of generating the observed symbol sequence given the assumed state sequence

Problem 1: Scoring and evaluation

- Given an observation sequence $O = \{ o_1, o_2, o_3, \dots o_T \}$ we want to compute the probability of generating it $P(O|\lambda)$
- We assume a sequence of states $q = \{ q_1, q_2, q_3, \dots q_T \}$
- Decompose the problem by summing over all possible state sequences:

$$P(O|\lambda) = \sum_{\text{all } q} P(O|q, \lambda) P(q|\lambda)$$

Likelihood of generating the observed symbol sequence given the assumed state sequence

How likely it is for the system to go through the given sequence of states

Problem 1: Scoring and evaluation

- Probability of the observation sequence given the state sequence:

$$P(O|q, \lambda) = \prod_{t=1}^T p(o_t | q_t, \lambda) = b_{q_1}(o_1) \cdot b_{q_2}(o_2) \cdots b_{q_T}(o_T)$$

- Probability of the state sequence:

$$P(q|\lambda) = \pi_{q_1} (a_{q_1 q_2}) \cdot (a_{q_2 q_3}) \cdots (a_{q_{T-1} q_T})$$

- Using the chain rule:

$$\begin{aligned} P(O|\lambda) &= \sum_{\text{all } q} P(O|q, \lambda) P(q|\lambda) \\ &= \sum_{\text{all } q} \pi_{q_1} b_{q_1}(o_1) a_{q_1 q_2} b_{q_2}(o_2) \cdots a_{q_{T-1} q_T} b_{q_T}(o_T) \end{aligned}$$


Not practical to compute!

Forward algorithm

- Define the probability of seeing observations \mathbf{o}_1 to \mathbf{o}_T , and ending in state i , given HMM λ :

$$\alpha_t(i) = P(\mathbf{o}_1 \mathbf{o}_2 \dots \mathbf{o}_t, q_t = i \mid \lambda)$$

Forward algorithm



Forward algorithm

- Define the probability of seeing observations \mathbf{o}_1 to \mathbf{o}_T , and ending in state i , given HMM λ :


$$\alpha_t(i) = P(\mathbf{o}_1 \mathbf{o}_2 \dots \mathbf{o}_t, q_t = i \mid \lambda)$$

- Initialization: $\alpha_0(i) = \pi_i$

- Induction:
$$\alpha_{t+1}(j) = \left[\sum_{i=1}^N \alpha_t(i) a_{ij} \right] b_j(\mathbf{o}_{t+1})$$


- Termination:
$$P(\mathbf{O} \mid \lambda) = \sum_{i=1}^N \alpha_T(i)$$

Forward algorithm example



- Given this HMM with discrete observations A and B, what is the probability of generating the sequence {A,A,B}?
- We need to calculate $P(\mathbf{O} = \{A,A,B\} \mid \lambda)$

Forward algorithm example




Backward algorithm

- Define the probability of seeing observations \mathbf{o}_{t+1} to \mathbf{o}_T , given state i at time t and HMM λ :

$$\beta_t(i) = P(\mathbf{o}_{t+1} \mathbf{o}_{t+2} \dots \mathbf{o}_T, q_t = i \mid \lambda)$$

Backward algorithm



Backward algorithm

- Define the probability of seeing observations \mathbf{o}_{t+1} to \mathbf{o}_T , given state i at time t and HMM λ :

$$\beta_t(i) = P(\mathbf{o}_{t+1} \mathbf{o}_{t+2} \dots \mathbf{o}_T, q_t = i \mid \lambda)$$


- Initialization: $\beta_T(i) = 1$

- Induction:
$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(\mathbf{o}_{t+1}) \beta_{t+1}(j)$$

$t = T-1, T-2, \dots, 1$
 $1 \leq i \leq N$

- Termination:
$$P(\mathbf{O} \mid \lambda) = \sum_{i=1}^N \pi_i \beta_0(i)$$

Backward algorithm example



Problem 2: Decoding

- Given an observation sequence $O = \{o_1, o_2, o_3, \dots, o_T\}$, and a model λ , how do we find the best sequence of states $q = \{q_1, q_2, q_3, \dots, q_T\}$ which maximizes $P(O, q | \lambda)$?
- Define the highest probable state sequence that accounts for observations o_1 to o_t and ends in state i at time t :

$$\delta_t(i) = \max_{q_1 q_2 \dots q_{t-1}} P(q_1 q_2 \dots q_{t-1}, q_t = i, o_1 o_2 \dots o_t | \lambda)$$

- At next transition:

$$\delta_{t+1}(j) = [\max_i \delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

Problem 1: Scoring and evaluation

- Solution: two ways of calculating $P(O | \lambda)$


- Forward algorithm

$$P(O | \lambda) = \sum_{i=1}^N \alpha_T(i)$$

- Backward algorithm

$$P(O | \lambda) = \sum_{i=1}^N \pi_i \beta_0(i)$$

Viterbi algorithm



Viterbi algorithm

Initialization:

$$\delta_1(i) = \pi_i b_i(\mathbf{o}_1)$$

$$\psi_1(i) = 0$$

Recursion:

$$\delta_t(j) = \max_{1 \leq i \leq N} [\delta_{t-1}(i) a_{ij}] b_j(\mathbf{o}_t)$$

$$\psi_t(j) = \arg \max_{1 \leq i \leq N} [\delta_{t-1}(i) a_{ij}]$$


Termination:

$$P^* = \max_{1 \leq i \leq N} [\delta_T(i)]$$

$$q_T^* = \arg \max_{1 \leq i \leq N} [\delta_T(i)]$$

Path back-tracing: $q_t^* = \psi_{t+1}(q_{t+1}^*)$

Viterbi algorithm example



Viterbi algorithm in log-domain

$$\tilde{\pi}_i = \log(\pi_i)$$

$$\tilde{b}_j(\mathbf{o}_t) = \log(b_j(\mathbf{o}_t))$$

$$\tilde{a}_{ij} = \log(a_{ij})$$

Same steps are followed in log-domain:

$$\delta_1(i) = \pi_i b_i(\mathbf{o}_1) \longrightarrow \tilde{\delta}_1(i) = \tilde{\pi}_i + \tilde{b}_i(\mathbf{o}_1)$$

Viterbi algorithm in log-domain

Initialization:

$$\tilde{\delta}_1(i) = \tilde{\pi}_i + \tilde{b}_i(\mathbf{o}_1)$$

$$\psi_1(i) = 0$$

Recursion:

$$\delta_t(j) = \max_{1 \leq i \leq N} [\tilde{\delta}_{t-1}(i) + \tilde{a}_{ij}] + \tilde{b}_j(\mathbf{o}_t)$$

$$\psi_t(j) = \arg \max_{1 \leq i \leq N} [\tilde{\delta}_{t-1}(i) + \tilde{a}_{ij}]$$

Termination:

$$\tilde{P}^* = \max_{1 \leq i \leq N} [\tilde{\delta}_T(i)]$$

$$q_T^* = \arg \max_{1 \leq i \leq N} [\tilde{\delta}_T(i)]$$

Path backtracking: $q_t^* = \psi_{t+1}(q_{t+1}^*)$


Problem 3: Training

- How do we tune λ to maximize $P(\mathbf{O} | \lambda)$?
 - No efficient algorithm to find global optimum
- Baum-Welch algorithm (forward-backward algorithm)
 - Iterative algorithm to find a local optimum
 - Compute probabilities using current model
 - Refine model parameters based on computed values

Forward-backward algorithm

- Define the probability of being in state i at time t and in state j at time $t+1$, given the model and the sequence


$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | \mathbf{O}, \lambda)$$



Forward-backward algorithm

- Define the probability of being in state i at time t and in state j at time $t+1$, given the model and the sequence

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(\mathbf{o}_{t+1}) \beta_{t+1}(j)}{P(\mathbf{O} | \lambda)} = \frac{\alpha_t(i) a_{ij} b_j(\mathbf{o}_{t+1}) \beta_{t+1}(j)}{\sum_{i=1}^N \alpha_t(i) \beta_t(i)}$$



Forward-backward algorithm

- More definitions, based on $\xi_t(i, j)$

$$\gamma_t(i) = \sum_{j=1}^N \xi_t(i, j)$$

Probability of being in state i at time t

$$\sum_{t=1}^{T-1} \gamma_t(i)$$

Expected number of transitions from state i in \mathbf{O}

$$\sum_{t=1}^{T-1} \sum_{j=1}^N \xi_t(i, j)$$

Expected number of transitions from state i to state j in \mathbf{O}

Computing the model parameters

- Initial state occupancy probability is the **expected number of times in state i at time $t=1$**

$$\bar{\pi}_i = \gamma_1(i)$$

Computing the model parameters

- transition probability from state i to state j

$$\bar{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

$$= \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

Computing the model parameters

- Probability of observing symbol k in state j

$$\bar{b}_j(k) = \frac{\text{expected number of times in state } j \text{ and observing } k\text{th symbol}}{\text{expected number times in state } j}$$

$$= \frac{\sum_{t=1}^T \gamma_t(j) \mathbf{1}_{o_t=v_k}}{\sum_{t=1}^T \gamma_t(j)} = \frac{\sum_{t=1}^T \alpha_t(j) \beta_t(j)}{\sum_{t=1}^T \alpha_t(j) \beta_t(j)}$$

Forward-backward algorithm iterations

1. Initialize $\lambda = (\mathbf{A}, \mathbf{B}, \pi)$
2. Compute α, β and ξ
3. Estimate $\bar{\lambda} = (\bar{\mathbf{A}}, \bar{\mathbf{B}}, \bar{\pi})$
4. Replace λ with $\bar{\lambda}$
5. Repeat from step 2 until convergence

Constraints:

$$\sum_{i=1}^N \bar{\pi}_i = 1$$

$$\sum_{j=1}^N \bar{a}_{ij} = 1 \quad 1 \leq i \leq N$$

$$\sum_{k=1}^M \bar{b}_j(k) = 1 \quad 1 \leq j \leq N$$

It can be shown that $P(\mathbf{O} | \bar{\lambda}) > P(\mathbf{O} | \lambda)$ unless $\bar{\lambda} = \lambda$

Mixture Gaussian PDFs

- Probability distribution of the state is a gaussian mixture

$$b_j(\mathbf{o}_t) = \sum_{k=1}^M c_{jk} \mathcal{N}(\mathbf{o}_t, \mu_{jk}, \Sigma_{jk}) \quad \sum_{k=1}^M c_{jk} = 1$$

$$c_{jk} \geq 0 \quad 1 \leq k \leq M$$

- Probability of being in state j at time t with the mixture component k accounting for the observation \mathbf{o}_t is:

$$\gamma_t(j, k) = \left[\frac{\alpha_t(j) \beta_t(j)}{\sum_{j=1}^N \alpha_t(j) \beta_t(j)} \right] \left[\frac{c_{jk} \mathcal{N}(\mathbf{o}_t, \mu_{jk}, \Sigma_{jk})}{\sum_{m=1}^M c_{jm} \mathcal{N}(\mathbf{o}_t, \mu_{jm}, \Sigma_{jm})} \right]$$

Parameter update equations for GMM PDF

- Mixture weight and mean:

$$\bar{c}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k)}{\sum_{t=1}^T \sum_{k'=1}^M \gamma_t(j, k')} \quad \bar{\mu}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k) \cdot \mathbf{o}_t}{\sum_{t=1}^T \gamma_t(j, k)}$$

- Transition matrix elements a_{ij} get updates same way as in the case of discrete symbols
- Covariance matrix:

$$\bar{\Sigma}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k) \cdot (\mathbf{o}_t - \bar{\mu}_{jk})(\mathbf{o}_t - \bar{\mu}_{jk})'}{\sum_{t=1}^T \gamma_t(j, k)}$$

Multiple observation sequences

- Variability in producing each sound unit is modeled by estimating HMM parameters from multiple examples of speech, collected from different speakers
- Assume K training observation sequences

$$\mathbf{O} = [\mathbf{O}^{(1)}, \mathbf{O}^{(2)}, \dots, \mathbf{O}^{(K)}]$$

$$\mathbf{O}^{(k)} = \{\mathbf{o}_1^k, \mathbf{o}_2^k, \dots, \mathbf{o}_{T_k}^k\}$$

$$P_k = P(\mathbf{O}^{(k)} | \lambda)$$

Parameter update for multiple observations

$$\bar{a}_{ij} = \frac{\sum_{k=1}^K \frac{1}{P_k} \sum_{t=1}^{T_k-1} \alpha_t^k(i) a_{ij} b_j(\mathbf{o}_{t+1}^{(k)}) \beta_{t+1}^k(j)}{\sum_{k=1}^K \frac{1}{P_k} \sum_{t=1}^{T_k-1} \alpha_t^k(i) \beta_t^k(i)}$$

$$\bar{b}_j(l) = \frac{\sum_{k=1}^K \frac{1}{P_k} \sum_{\substack{t=1 \\ s.t. \mathbf{o}_t = \mathbf{o}_l}}^{T_k-1} \alpha_t^k(i) \beta_t^k(j)}{\sum_{k=1}^K \frac{1}{P_k} \sum_{t=1}^{T_k-1} \alpha_t^k(i) \beta_t^k(i)}$$

One-state HMM with M component GMM

- Observation probability is a GMM with M components

$$b(\mathbf{o}_t) = \sum_{k=1}^M w_k b_k(\mathbf{o}_t, \mu_k, \Sigma_k)$$

- Probability that \mathbf{o}_t is generated by the k th component

$$P(k | \mathbf{o}_t, \lambda) = \frac{w_k b_k(\mathbf{o}_t)}{\sum_{k=1}^M w_k b_k(\mathbf{o}_t)}$$

Update equations for one-state HMM

Updated (new) parameter estimates

$$\begin{aligned} \bar{w}_k &= \frac{1}{T} \sum_{t=1}^T P(k | \mathbf{o}_t, \lambda) \\ \bar{\mu}_k &= \frac{\sum_{t=1}^T P(k | \mathbf{o}_t, \lambda) \cdot \mathbf{o}_t}{\sum_{t=1}^T P(k | \mathbf{o}_t, \lambda)} \\ \bar{\Sigma}_k &= \frac{\sum_{t=1}^T P(k | \mathbf{o}_t, \lambda) \cdot (\mathbf{o}_t)^2}{\sum_{t=1}^T P(k | \mathbf{o}_t, \lambda)} - (\bar{\mu}_k)^2 \end{aligned}$$

This term is computed using model parameters from previous algorithm iteration.